

```
In [1]: import pandas as pd
import numpy as np
import matplotlib.pyplot as plt
import math
import random
import seaborn as sns

%matplotlib inline

from IPython.display import Image
```

```
In [2]: heart_df = pd.read_csv('./DataSet/heart.csv')
```

```
In [3]: heart_df.head()
```

```
Out[3]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
0	63	1	3	145	233	1	0	150	0	2.3	0	0	1	1
1	37	1	2	130	250	0	1	187	0	3.5	0	0	2	1
2	41	0	1	130	204	0	0	172	0	1.4	2	0	2	1
3	56	1	1	120	236	0	1	178	0	0.8	2	0	2	1
4	57	0	0	120	354	0	1	163	1	0.6	2	0	2	1

```
In [4]: heart_df.tail()
```

```
Out[4]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang	oldpeak	slope	ca	thal	target
298	57	0	0	140	241	0	1	123	1	0.2	1	0	3	0
299	45	1	3	110	264	0	1	132	0	1.2	1	0	3	0
300	68	1	0	144	193	1	1	141	0	3.4	1	2	3	0
301	57	1	0	130	131	0	1	115	1	1.2	1	1	3	0
302	57	0	1	130	236	0	0	174	0	0.0	1	1	2	0

```
In [5]: heart_df.shape
```

```
Out[5]: (303, 14)
```

```
In [6]: heart_df.info()
```

```
<class 'pandas.core.frame.DataFrame'>
RangeIndex: 303 entries, 0 to 302
Data columns (total 14 columns):
 #   Column      Non-Null Count  Dtype
---  -
 0   age         303 non-null    int64
 1   sex         303 non-null    int64
 2   cp          303 non-null    int64
 3   trestbps    303 non-null    int64
 4   chol        303 non-null    int64
```

```

5   fbs          303 non-null    int64
6   restecg     303 non-null    int64
7   thalach     303 non-null    int64
8   exang       303 non-null    int64
9   oldpeak     303 non-null    float64
10  slope       303 non-null    int64
11  ca          303 non-null    int64
12  thal        303 non-null    int64
13  target      303 non-null    int64
dtypes: float64(1), int64(13)
memory usage: 33.3 KB

```

```
In [7]: heart_df.isnull().sum()
```

```

Out[7]: age          0
sex          0
cp           0
trestbps     0
chol         0
fbs          0
restecg      0
thalach      0
exang        0
oldpeak      0
slope        0
ca           0
thal         0
target       0
dtype: int64

```

```
In [8]: heart_df.describe()
```

```

Out[8]:
```

	age	sex	cp	trestbps	chol	fbs	restecg	thalach	exang
<b>count</b>	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000	303.000000
<b>mean</b>	54.366337	0.683168	0.966997	131.623762	246.264026	0.148515	0.528053	149.646865	0.326733
<b>std</b>	9.082101	0.466011	1.032052	17.538143	51.830751	0.356198	0.525860	22.905161	0.469794
<b>min</b>	29.000000	0.000000	0.000000	94.000000	126.000000	0.000000	0.000000	71.000000	0.000000
<b>25%</b>	47.500000	0.000000	0.000000	120.000000	211.000000	0.000000	0.000000	133.500000	0.000000
<b>50%</b>	55.000000	1.000000	1.000000	130.000000	240.000000	0.000000	1.000000	153.000000	0.000000
<b>75%</b>	61.000000	1.000000	2.000000	140.000000	274.500000	0.000000	1.000000	166.000000	1.000000
<b>max</b>	77.000000	1.000000	3.000000	200.000000	564.000000	1.000000	2.000000	202.000000	1.000000

```
In [9]: heart_df.columns
```

```

Out[9]: Index(['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach',
              'exang', 'oldpeak', 'slope', 'ca', 'thal', 'target'],
              dtype='object')

```

```

In [10]: X = heart_df[['age', 'sex', 'cp', 'trestbps', 'chol', 'fbs', 'restecg', 'thalach',
                      'exang', 'oldpeak', 'slope', 'ca', 'thal']]
          Y = heart_df['target']

```

```

In [11]: Y = np.array((Y - Y.mean())/Y.std())
          X = X.apply(lambda rec:(rec-rec.mean())/rec.std(),axis=0) #axis=0 -> It is applied to all

```

```
print(X)
```

```
[303 rows x 13 columns]
```

```
print(Y)
```

[illegible]

[illegible]

## Working on the Gradient Descent Algorithm

```
In [14]: Image(url= "Image/StepsForGradientDescent.png")
```

Out[14]: 

```
In [15]: def initialize(dim):
          np.random.seed(seed=42)
          random.seed(42)
          b = random.random()           # b -> beta0 (bias or intercept)
          w = np.random.rand(dim)       # w -> beta1 (feature weight)
          return b,w
```

```
In [16]: def predict_Y(b,w,x):  
          return b + np.matmul(x,w)
```

```
In [17]: def get_cost(Y,Y_hat):
          Y_resid = Y - Y_hat
          return np.sum(np.matmul(Y_resid.T, Y_resid))/len(Y_resid)
```

```
In [18]: def update_beta(x, y, y_hat, b_0, w_0, learning_rate):
          db = (np.sum(y_hat - y)*2)/len(y)
          dw = (np.dot(y_hat - y, x)*2)/len(y)
          b_1 = b_0 - learning_rate*db
          w_1 = w_0 - learning_rate*dw
          return b_1, w_1
```

```
In [19]: def run_gradient_descent(X, Y, alpha=0.01, num_iterations=100):
          b,w = initialize(X.shape[1])

          iter_num = 0
          gd_iteration_df = pd.DataFrame(columns=['Iterations','Cost'])
          result_idx = 0

          for each_iter in range(num_iterations):
              y_hat = predict_Y(b, w, X)
```

```

    this_cost = get_cost(Y, y_hat)
    prev_b = b
    prev_w = w
    b, w = update_beta(X, Y, y_hat, prev_b, prev_w, alpha)

    if iter_num%10==0:
        gd_iteration_df.loc[result_idx] = [iter_num, this_cost]
        result_idx = result_idx + 1
        iter_num = iter_num + 1

print("\nFinal Estimate of b is {} \n w is {}".format(b,w))
return gd_iteration_df, b, w

```

In [20]: `gid, b , w = run_gradient_descent(X, Y, alpha=0.001, num_iterations=200)`

```

Final Estimate of b is 0.42844895817391493
w is [0.17489273 0.49510225 0.70118669 0.27438212 0.06249845 0.00677069
0.13890371 0.85186318 0.31733192 0.29896375 0.18352353 0.48662455
0.3574579 ]

```


In [21]: `gid.head(20)`

Out[21]:

	Iterations	Cost
0	0.0	9.376909
1	10.0	8.785926
2	20.0	8.243039
3	30.0	7.743802
4	40.0	7.284208
5	50.0	6.860642
6	60.0	6.469848
7	70.0	6.108885
8	80.0	5.775100
9	90.0	5.466095
10	100.0	5.179706
11	110.0	4.913976
12	120.0	4.667134
13	130.0	4.437577
14	140.0	4.223855
15	150.0	4.024652
16	160.0	3.838778
17	170.0	3.665151
18	180.0	3.502789
19	190.0	3.350799

In [22]: `plt.plot(gid['Iterations'],gid['Cost'])`  
`plt.xlabel('Number of Iterations')`

```
plt.ylabel('Cost of MSE')
plt.show()
```



In [23]:

```
alpha_df_1, b, w = run_gradient_descent(X, Y, alpha=0.1, num_iterations=2000)
alpha_df_2, b, w = run_gradient_descent(X, Y, alpha=0.01, num_iterations=2000)
alpha_df_3, b, w = run_gradient_descent(X, Y, alpha=0.001, num_iterations=2000)
alpha_df_4, b, w = run_gradient_descent(X, Y, alpha=0.0001, num_iterations=2000)
```

Final Estimate of b is -6.271102523981385e-17

```
w is [-0.014937 -0.1830988 0.23317503 -0.0699988 -0.03673396 0.01240582
0.05254851 0.13863953 -0.13566011 -0.13683504 0.09756498 -0.20623343
-0.14611027]
```

Final Estimate of b is -5.998397015702797e-17

```
w is [-0.01493697 -0.18309879 0.23317503 -0.06999881 -0.03673397 0.01240581
0.05254851 0.13863957 -0.13566009 -0.13683508 0.09756494 -0.20623342
-0.14611027]
```

Final Estimate of b is 0.011664695556930303

```
w is [ 0.03566807 -0.17329831 0.27600757 -0.08635701 -0.04145347 -0.01574182
0.06678565 0.23549895 -0.05319714 -0.11438976 0.09934815 -0.17545142
-0.1648434 ]
```

Final Estimate of b is 0.4286034542113327

```
w is [0.17517122 0.49551784 0.70107311 0.27472357 0.06260663 0.00696353
0.13876063 0.85166111 0.31767097 0.29944636 0.18323055 0.48711162
0.35795044]
```

In [24]:

```
plt.plot(alpha_df_1['Iterations'],alpha_df_1['Cost'],label='alpha=0.1')
plt.plot(alpha_df_2['Iterations'],alpha_df_2['Cost'],label='alpha=0.01')
plt.plot(alpha_df_3['Iterations'],alpha_df_3['Cost'],label='alpha=0.001')
plt.plot(alpha_df_4['Iterations'],alpha_df_4['Cost'],label='alpha=0.0001')
plt.legend()
plt.xlabel('Number of Iterations')
plt.ylabel('Cost')
plt.title('Heart Disease Prediction using Gradient Descent Algorithm')
```

Out[24]:

Text(0.5, 1.0, 'Heart Disease Prediction using Gradient Descent Algorithm')

Heart Disease Prediction using Gradient Descent Algorithm

